




Edge Detection




- Edges: **Boundary** (steep-changes) between two regions with distinct gray-level properties.
- Most important visual information.

How can we detect changes?




Edge Detection – Derivatives!


- Steepness is defined by spatial derivatives: $\frac{\partial I}{\partial x}$ and $\frac{\partial I}{\partial y}$ (and $\frac{\partial I}{\partial z}$ in 3D)
- Image edges (gradients) have direction and magnitude



Forward difference




Central difference



Diagonal difference


Derivatives

Reminder





Derivatives and edges

An edge is a place of **rapid change** in the image intensity function.



Edge Detection

Using derivatives



Derivatives with convolution

Where would the function “go” if we give it a little “nudge”?

For 2D function, $f(x, y)$, the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\epsilon \rightarrow 0} \frac{f(x + \epsilon, y) - f(x, y)}{\epsilon}$$

For discrete data, we can approximate using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

Edge Detection as Derivatives

Where would the function “go” if we give it a little “nudge”?

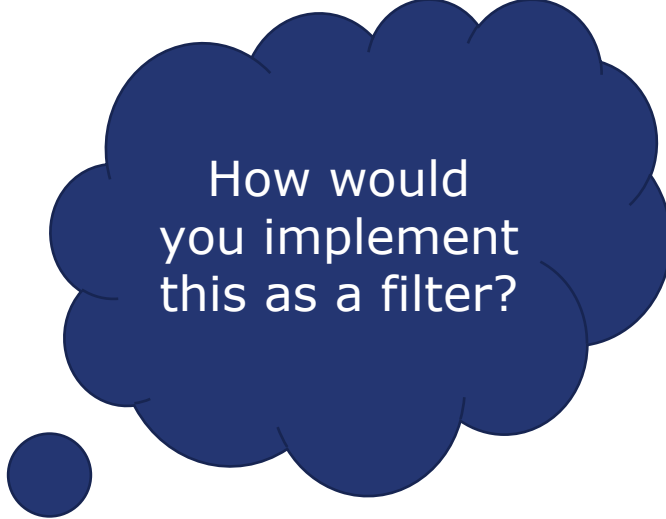
- For continuous functions: $f'(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$

- Approximate partial derivatives for both x and y:

$$\frac{\partial}{\partial x} f(i, j) \cong f(i, j) - f(i - 1, j)$$

$$\frac{\partial}{\partial y} f(i, j) \cong f(i, j) - f(i, j - 1)$$

- Approximate derivatives by convolution with $\begin{bmatrix} 1 & -1 \end{bmatrix}$




How would you implement this as a filter?


Edge Detection Convolution Kernels

Calculating Gradients


- Forward Finite Difference: $f'(x) = \lim_{\Delta x \rightarrow 0} \frac{f(x+\Delta x) - f(x)}{\Delta x}$



- Backward Finite Difference: $f'(x) = \lim_{\Delta x \rightarrow 0} \frac{f(x) - f(x-\Delta x)}{\Delta x}$



- Central Finite Difference: $f'(x) = \lim_{\Delta x \rightarrow 0} \frac{f(x+\Delta x) - f(x-\Delta x)}{2\Delta x}$



These kernels perform well on very clean images. They don't work as well if there's noise on the images (kernels are too small!)

Derivatives Approximation with Kernel Operations

- Roberts Operator
- Prewitt Operator
- Sobel Operator

+1	0	-1
+1	0	-1
+1	0	-1

Gx

-1	-1	-1
0	0	0
+1	+1	+1

Gy

Prewitt Operator

-1	0	+1
-2	0	+2
-1	0	+1

Gx

+1	+2	+1
0	0	0
-1	-2	-1

Gy

Sobel Operator




Gradient Operators

- 2D first derivatives - $\frac{\partial}{\partial x} f(i, j) \cong f(i, j) - f(i - 1, j)$
- $\nabla f = \begin{bmatrix} G_x \\ G_y \end{bmatrix} = \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix}$
- Gradient Magnitude: $\nabla f = \sqrt{G_x^2 + G_y^2} \approx |G_x| + |G_y|$
- Gradient Direction: $\alpha(x, y) = \tan^{-1} \left(\frac{G_y}{G_x} \right)$

Image Gradients

The gradient of an image: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$

The gradient points in the direction of most rapid change in intensity




The **gradient direction** (orientation of edge normal) is given by:

$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$


The **edge strength** is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x} \right)^2 + \left(\frac{\partial f}{\partial y} \right)^2}$$



Edge Detection

Using derivatives for a discrete function



Gradient vector:


$$\nabla I = \left[\frac{\partial I}{\partial x}, \frac{\partial I}{\partial y} \right]^T$$

Magnitude:

$$|\nabla I| = \sqrt{\left(\frac{\partial I}{\partial x} \right)^2 + \left(\frac{\partial I}{\partial y} \right)^2}$$


Orientation:

$$\theta = \text{atan}^2 \left(\frac{\partial I}{\partial y}, \frac{\partial I}{\partial x} \right)$$



Edge Detection

Using derivatives for a discrete function through filters



+1	0	-1
+1	0	-1
+1	0	-1



Gx

-1	-1	-1
0	0	0
+1	+1	+1


Gy

Edge Detection

Using derivatives for a discrete function through filters convolutions



Sobel Operator



- Approximate gradient calculation (the image derivatives)

- The final edge is received by L2-Distance:

$$\sqrt{G_x^2 + G_y^2}$$

-1	0	+1
-2	0	+2
-1	0	+1

Gx

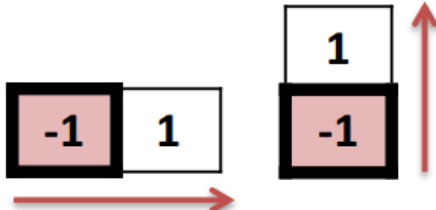
+1	+2	+1
0	0	0
-1	-2	-1

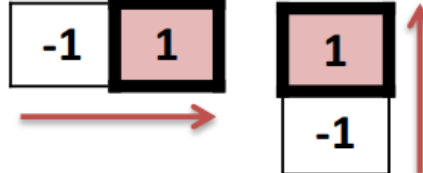
Gy

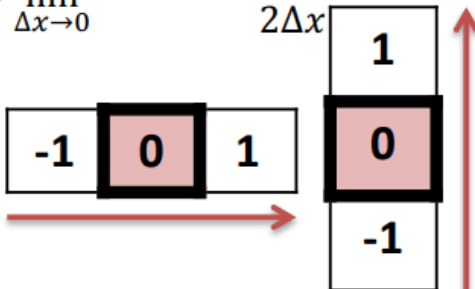
Edge Detection Convolution Kernels

Prewitt and Sobel compute derivatives in one direction while smoothing in the other direction

Calculating Gradients

$$f'(x) = \lim_{\Delta x \rightarrow 0} \frac{f(x + \Delta x) - f(x)}{\Delta x}$$


$$f'(x) = \lim_{\Delta x \rightarrow 0} \frac{f(x) - f(x - \Delta x)}{\Delta x}$$


$$f'(x) = \lim_{\Delta x \rightarrow 0} \frac{f(x + \Delta x) - f(x - \Delta x)}{2\Delta x}$$


Forward Finite Difference

Backward Finite Difference

Central Finite Difference

-1	0	1
-1	0	1
-1	0	1

1	1	1
0	0	0
-1	-1	-1


Prewitt Operator

-1	0	1
-2	0	2
-1	0	1

1	2	1
0	0	0
-1	-2	-1

Sobel Operator

Sobel Operator



CANNY EDGE DETECTION

When noise is on the way


Canny Edge Detector

Classic Algorithm for Finding Edges in Images

- We want a kernel to simultaneously maximize SNR and localization for a step edge under white *Gaussian* noise
- Solution by numerical optimization is very close to derivative of Gaussian kernel

Numerically Optimized


Derivative of Gaussian



Canny, IEEE TPAMI 1986

Edge Detection - Laplacian

Second Derivative



- Laplacians: sum of second derivatives:

$$\nabla^2 I = \frac{\partial^2 I}{\partial x^2} + \frac{\partial^2 I}{\partial y^2}$$


- Zero-crossings mark edge location
- Do not provide directions of edges

Edge Detection - Laplacian

Second Derivative – the difference of the difference

- $\frac{\partial^2 I}{\partial x^2} \approx \frac{1}{\epsilon^2} (I_{i-1,j} - 2I_{i,j} + I_{i+1,j})$
- $\frac{\partial^2 I}{\partial y^2} \approx \frac{1}{\epsilon^2} (I_{i,j-1} - 2I_{i,j} + I_{i,j+1})$
- $\nabla^2 I = \frac{\partial^2 I}{\partial x^2} + \frac{\partial^2 I}{\partial y^2}$




$$\nabla^2 = \frac{1}{\epsilon^2} \begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix} \text{ or (better) } \nabla^2 = \frac{1}{6\epsilon^2} \begin{bmatrix} 1 & 4 & 1 \\ 4 & -20 & 4 \\ 1 & 4 & 1 \end{bmatrix}$$



Edge Detection - Laplacian


Laplacian of a gaussian

0	0	3	2	2	2	3	0	0
0	2	3	5	5	5	3	2	0
3	3	5	3	0	3	5	3	3
2	5	3	-12	-23	-12	3	5	2
2	5	0	-23	-40	-23	0	5	2
2	5	3	-12	-23	-12	3	5	2



$\sqrt{2}$

$$Log(x, y) = -\frac{1}{\pi\sigma^4} \left(1 - \frac{x^2 + y^2}{2\sigma^2} \right) e^{-\frac{x^2 + y^2}{2\sigma^2}}$$




(Partial) Derivative of Gaussian Kernel

1D



$$G(\cdot)$$



$$\frac{\partial G(\cdot)}{\partial x}$$



2D



3D

$$\frac{1}{(\sigma\sqrt{2\pi})^3} e^{-\frac{(x^2+y^2+z^2)}{2\sigma^2}}$$

$$-\frac{x}{\sigma^2(\sigma\sqrt{2\pi})^3} e^{-\frac{(x^2+y^2+z^2)}{2\sigma^2}}$$


A Computational Approach to Edge Detection

John Canny

Canny Edge Detection

Commonly used edge-detection filter

- Noise smoothing with a gaussian filter $\rightarrow n_{\sigma} * I$
- Edge gradients (e.g., with Sobel operator) $\rightarrow \nabla n_{\sigma} * I$
- Find Gradient Magnitude for each pixel $\rightarrow \|\nabla n_{\sigma} * I\|$
- Find Gradient Orientation for each pixel $\rightarrow \hat{n} = \frac{\nabla n_{\sigma} * I}{\|\nabla n_{\sigma} * I\|}$
- Compute the Laplacian along the gradient direction \hat{n} at each pixel $\rightarrow \frac{\partial^2 n_{\sigma} * I}{\partial \hat{n}^2}$



Canny Edge Detection

In other words

- Convoluting with a derivative of Gaussian is the same as convoluting with a Gaussian and then taking a derivative
(...however, taking the *exact* derivative of an image is not well defined)
- It is (more or less) equivalent to blurring first, to reduce noise, then taking an exact derivative.

Canny Edge Detection

Final steps:

- Hysteresis thresholding:
 - First, global threshold with an upper threshold (on $|\nabla I|$)
 - Then, region grow starting from those regions using a lower threshold
- Non-maximum suppression
 - Eliminate edges that are not local maxima (along the direction of the gradient)

Hysteresis thresholding


- Threshold at low/high levels to get weak/strong edge pixels
- Do connected components, starting from strong edge pixels



Credit: James Hays


Hysteresis thresholding

- Use a high threshold to start edge curves, and a low threshold to continue them.



Canny Edge Detection

The filtering's sigma value affects the results



Take-Aways

- There are many different segmentation methods
 - Thresholding
 - Region Growing
 - ML-Based
 - And even more is coming up soon:
 - Level-sets,
 - Fast-Marching,
 - Graph-based
- Edge Detection
 - Can produce disconnected boundaries
 - Over/Under Segmentation
 - Harder in 3D
 - Canny smooths before applying gradients
- No one algorithm is the “best”.
 - Must match to the problem and data